Compression and Contraction

Adway Girish
Information Theory Lab

EPFL Information Processing Group

September 9, 2024
IPG PhD Review

Outline

1. Prompt compression for black-box language models
2. Input-entropy-constrained capacity
3. Joint range of divergences
4. Closing remarks

Outline

Prompt compression

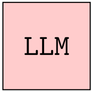

Prompt compression

Prompt

It was the best of times, it was the worst of times, it was the age of wisdom, it was the age of foolishness, it was the epoch of belief, it was the epoch of incredulity, it was the season of light, it was the season of darkness, it was the spring of hope, it was the winter of despair.

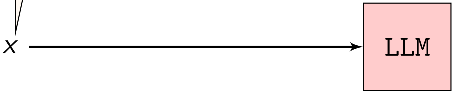

Prompt compression

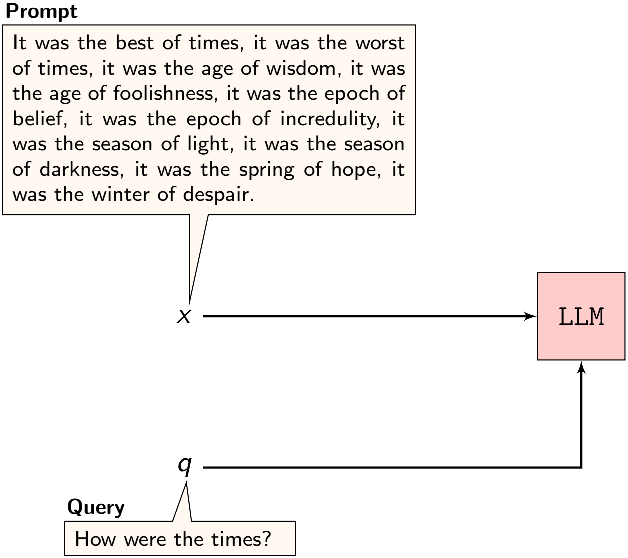

Prompt compression

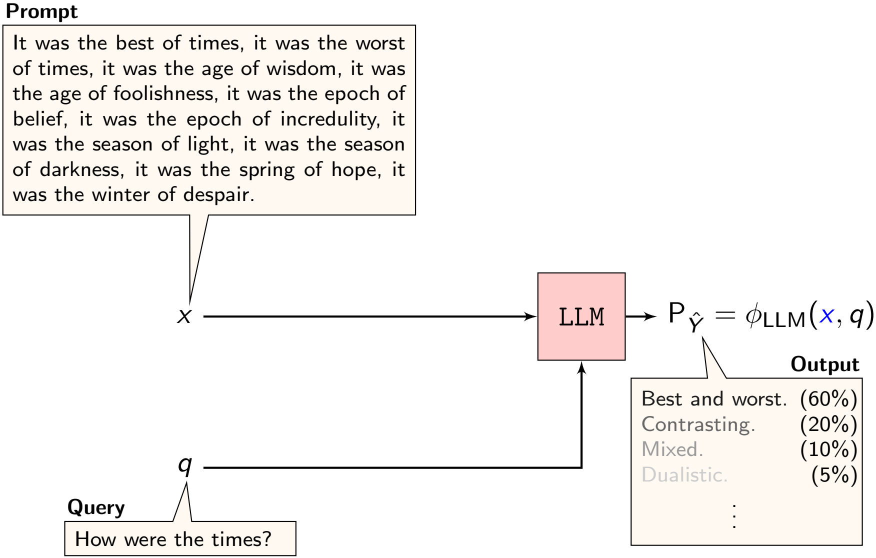

Prompt compression: query-agnostic

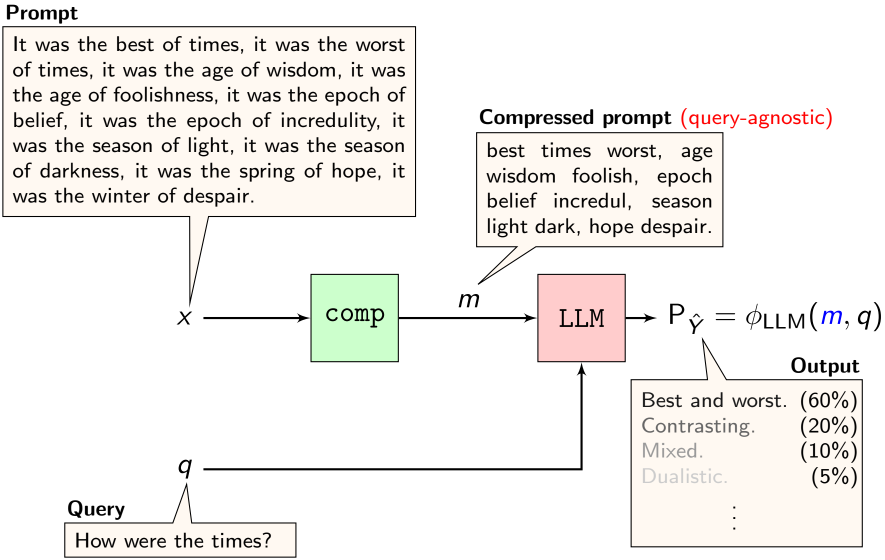

Prompt compression: query-aware

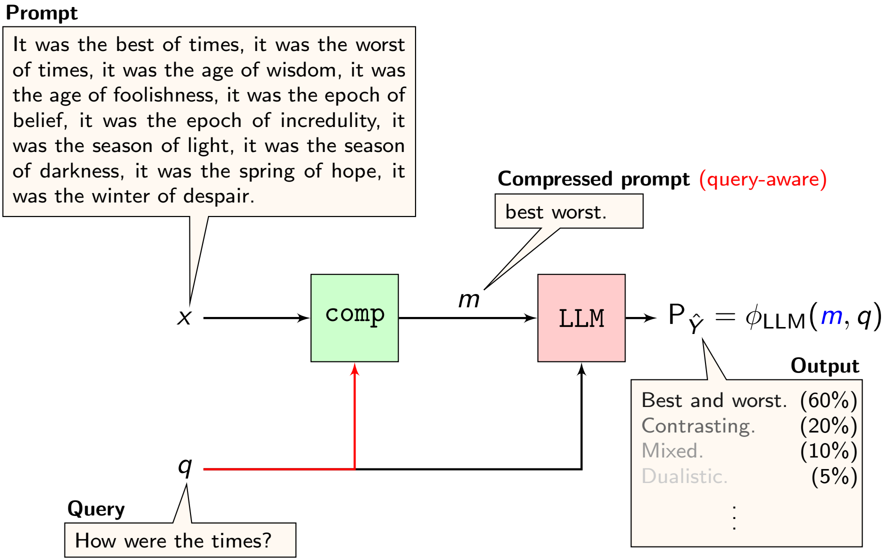

Prompt compression: rate-distortion formulation

- $(X, Q, Y) \sim \mathrm{P}_{XQY} = \mathrm{P}_{XQ}\,\mathrm{P}_{Y|XQ}$ $\qquad$ $Y =$ "true answer"

Prompt compression: rate-distortion formulation

- $(X, Q, Y) \sim \mathsf{P}_{XQY} = \mathsf{P}_{XQ}\,\mathsf{P}_{Y|XQ}$ $\quad Y =$ "true answer"

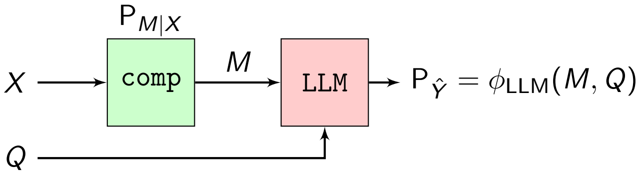

Prompt compression: rate-distortion formulation

- $(X, Q, Y) \sim \mathsf{P}_{XQY} = \mathsf{P}_{XQ}\,\mathsf{P}_{Y|XQ}$ $\qquad$ $Y =$ "true answer"

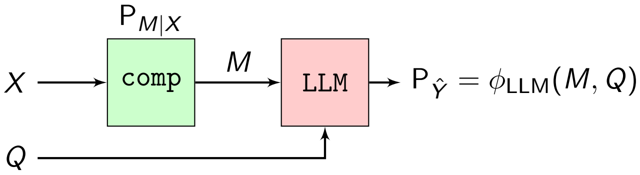

- Compression with side-information

Prompt compression: rate-distortion formulation

- $(X, Q, Y) \sim \mathrm{P}_{XQY} = \mathrm{P}_{XQ}\,\mathrm{P}_{Y|XQ}$ $\qquad Y =$ "true answer"

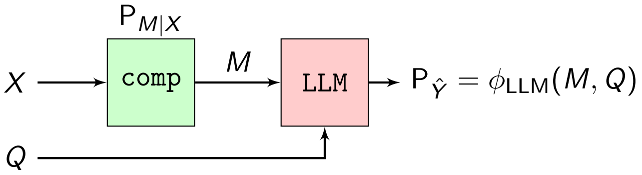

- Compression with side-information
  for a fixed decoder, "$(m, q) \mapsto \phi_{\mathrm{LLM}}(m, q)$"

Prompt compression: rate-distortion formulation

- $(X, Q, Y) \sim \mathrm{P}_{XQY} = \mathrm{P}_{XQ}\,\mathrm{P}_{Y|XQ}$ $\qquad Y =$ "true answer"

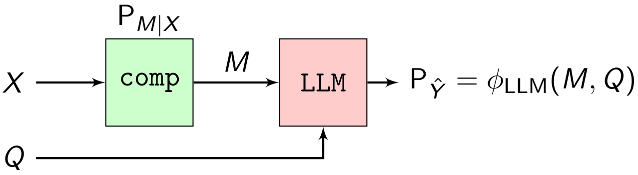

- Compression with side-information
  for a fixed decoder, "$(m, q) \mapsto \phi_{\mathrm{LLM}}(m, q)$"

- Performance metrics:

Prompt compression: rate-distortion formulation

- $(X, Q, Y) \sim \mathsf{P}_{XQY} = \mathsf{P}_{XQ}\,\mathsf{P}_{Y|XQ}$ $\qquad Y =$ "true answer"

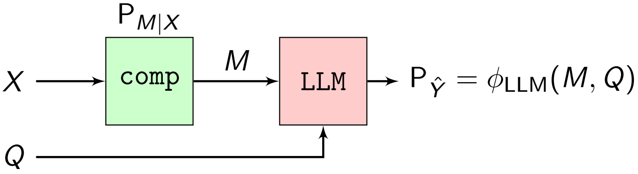

- Compression with side-information

  for a fixed decoder, "$(m, q) \mapsto \phi_{\mathsf{LLM}}(m, q)$"

- Performance metrics:

$$\text{rate} = \mathbb{E}\left[\frac{\mathrm{len}(M)}{\mathrm{len}(X)}\right] \qquad \text{distortion} = \mathbb{E}\left[\mathrm{d}\big(Y, \phi_{\mathsf{LLM}}(M, Q)\big)\right]$$

Distortion-rate function

- $$\text{rate} = \mathbb{E}\left[\frac{\text{len}(M)}{\text{len}(X)}\right] \qquad \text{distortion} = \mathbb{E}\left[\text{d}(Y, \phi_{\text{LLM}}(M, Q))\right]$$

Distortion-rate function

- $$\text{rate} = \mathbb{E}\left[\frac{\text{len}(M)}{\text{len}(X)}\right] \qquad \text{distortion} = \mathbb{E}\left[\text{d}(Y, \phi_{\text{LLM}}(M, Q))\right]$$
- $$\begin{aligned} D^*(R) = \quad & \inf_{\text{P}_{M|X}} && \mathbb{E}\left[\text{d}\big(Y, \phi_{\text{LLM}}(M, Q)\big)\right] \\ & \text{s.t.} && \mathbb{E}\left[\frac{\text{len}(M)}{\text{len}(X)}\right] \leq R, \text{ and} \\ & && \text{P}_{M|X} \text{ ``is a compressor''} \end{aligned}$$

Distortion-rate function

- $$\text{rate} = \mathbb{E}\left[\frac{\text{len}(M)}{\text{len}(X)}\right] \qquad \text{distortion} = \mathbb{E}\left[\text{d}(Y, \phi_{\text{LLM}}(M, Q))\right]$$
- $$\begin{aligned} D^*(R) = \quad \inf_{\text{P}_{M|X}} \quad & \mathbb{E}\left[\text{d}\big(Y, \phi_{\text{LLM}}(M, Q)\big)\right] \\ \text{s.t.} \quad & \mathbb{E}\left[\frac{\text{len}(M)}{\text{len}(X)}\right] \le R, \text{ and} \\ & \text{P}_{M|X} \text{ ``is a compressor''} \end{aligned}$$
- Linear program, but large dimension

Distortion-rate function

- $$\text{rate} = \mathbb{E}\left[\frac{\text{len}(M)}{\text{len}(X)}\right] \qquad \text{distortion} = \mathbb{E}\left[\text{d}(Y, \phi_{\text{LLM}}(M, Q))\right]$$

- $$\begin{aligned} D^*(R) = \quad \inf_{\mathrm{P}_{M|X}} \quad & \mathbb{E}\left[\text{d}\big(Y, \phi_{\text{LLM}}(M, Q)\big)\right] \\ \text{s.t.} \quad & \mathbb{E}\left[\frac{\text{len}(M)}{\text{len}(X)}\right] \le R, \text{ and} \\ & \mathrm{P}_{M|X} \text{ ``is a compressor''} \end{aligned}$$

- Linear program, but large dimension $\approx 32{,}000^{10}$

Distortion-rate function

- $$\text{rate} = \mathbb{E}\left[\frac{\text{len}(M)}{\text{len}(X)}\right] \qquad \text{distortion} = \mathbb{E}\left[\text{d}(Y, \phi_{\text{LLM}}(M, Q))\right]$$

- $$\begin{aligned} D^*(R) = \quad \inf_{\mathsf{P}_{M|X}} \quad & \mathbb{E}\left[\text{d}\big(Y, \phi_{\text{LLM}}(M, Q)\big)\right] \\ \text{s.t.} \quad & \mathbb{E}\left[\frac{\text{len}(M)}{\text{len}(X)}\right] \leq R, \text{ and} \\ & \mathsf{P}_{M|X} \text{ ``is a compressor''} \end{aligned}$$

- Linear program, but large dimension $\approx 32{,}000^{10}$

- Dual:
$$D^*(R) \;=\; \sup_{\lambda \geq 0} \left\{ -\lambda R + \sum_{x \in \mathcal{X}} \min_{m \in \mathcal{M}_x} \left[\boldsymbol{D}_{x,m} + \lambda \, \boldsymbol{R}_{x,m} \right] \right\}$$

Distortion-rate function

- $$\text{rate} = \mathbb{E}\left[\frac{\text{len}(M)}{\text{len}(X)}\right] \qquad \text{distortion} = \mathbb{E}\left[\text{d}(Y, \phi_{\text{LLM}}(M, Q))\right]$$

- $$\begin{aligned} D^*(R) = \quad & \inf_{\mathsf{P}_{M|X}} \quad \mathbb{E}\left[\text{d}\left(Y, \phi_{\text{LLM}}(M, Q)\right)\right] \\ & \text{s.t.} \quad \mathbb{E}\left[\frac{\text{len}(M)}{\text{len}(X)}\right] \leq R, \text{ and} \\ & \qquad \mathsf{P}_{M|X} \text{ “is a compressor”} \end{aligned}$$

- Linear program, but large dimension $\approx 32{,}000^{10}$

- Dual:
$$D^*(R) = \sup_{\lambda \geq 0} \left\{ -\lambda R + \sum_{x \in \mathcal{X}} \min_{m \in \underset{\text{all possible “compressions” of } x}{\mathcal{M}_x}} \left[\boldsymbol{D}_{x,m} + \lambda \boldsymbol{R}_{x,m} \right] \right\}$$

Distortion-rate function

- $$\text{rate} = \mathbb{E}\left[\frac{\text{len}(M)}{\text{len}(X)}\right] \qquad \text{distortion} = \mathbb{E}\left[\text{d}(Y, \phi_{\text{LLM}}(M, Q))\right]$$

- $$\begin{aligned} D^*(R) = \quad \inf_{\mathrm{P}_{M|X}} \quad & \mathbb{E}\left[\text{d}\big(Y, \phi_{\text{LLM}}(M, Q)\big)\right] \\ \text{s.t.} \quad & \mathbb{E}\left[\frac{\text{len}(M)}{\text{len}(X)}\right] \le R, \text{ and} \\ & \mathrm{P}_{M|X} \text{ ``is a compressor''} \end{aligned}$$

- Linear program, but large dimension $\approx 32{,}000^{10}$

- Dual:

$$D^*(R) \; = \; \sup_{\lambda \ge 0} \left\{ -\lambda R + \sum_{x \in \mathcal{X}} \min_{m \in \overbrace{\mathcal{M}_x}^{\text{all possible ``compressions'' of } x}} \left[ \underbrace{\boldsymbol{D}_{x,m} + \lambda \boldsymbol{R}_{x,m}}_{\text{``normalized'' distortion, rate on compressing } x \;\mapsto\; m} \right] \right\}$$

Distortion-rate function: geometric solution via dual

- Dual: $$D^*(R) \;=\; \sup_{\lambda \geq 0} \left\{ -\lambda R + \sum_{x \in \mathcal{X}} \min_{m \in \mathcal{M}_x} [\boldsymbol{D}_{x,m} + \lambda \boldsymbol{R}_{x,m}] \right\}$$

Distortion-rate function: geometric solution via dual

- Dual: $$D^*(R) \;=\; \sup_{\lambda \geq 0} \left\{ -\lambda R + \sum_{x \in \mathcal{X}} \underbrace{\min_{m \in \mathcal{M}_x} \left[\boldsymbol{D}_{x,m} + \lambda \boldsymbol{R}_{x,m} \right]}_{\text{for fixed } (\lambda, x)} \right\}$$

Distortion-rate function: geometric solution via dual

- Dual: $$D^*(R) = \sup_{\lambda \geq 0} \Big\{ -\lambda R + \sum_{x \in \mathcal{X}} \underbrace{\min_{m \in \mathcal{M}_x} [\boldsymbol{D}_{x,m} + \lambda \boldsymbol{R}_{x,m}]}_{\text{for fixed } (\lambda, x)} \Big\}$$

- Fix $\lambda \geq 0$, $x \in \mathcal{X}$

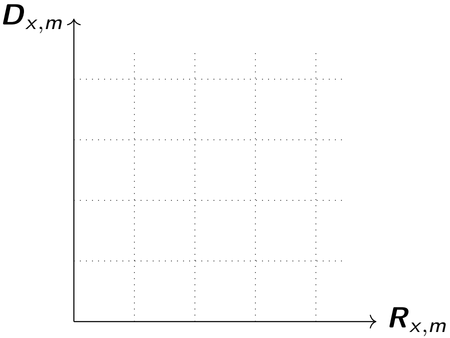

Distortion-rate function: geometric solution via dual

- Dual: $$D^*(R) \;=\; \sup_{\lambda \geq 0} \left\{ -\lambda R + \sum_{x \in \mathcal{X}} \underbrace{\min_{m \in \mathcal{M}_x} \left[\boldsymbol{D}_{x,m} + \lambda \boldsymbol{R}_{x,m} \right]}_{\text{for fixed } (\lambda, x)} \right\}$$

- Fix $\lambda \geq 0$, $x \in \mathcal{X}$

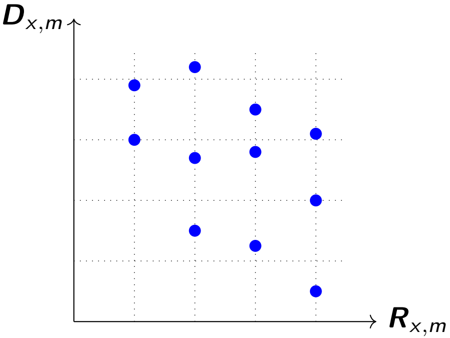

- Relevant points: $32{,}000^{10}$

Distortion-rate function: geometric solution via dual

- Dual: $$D^*(R) \;=\; \sup_{\lambda \geq 0} \left\{ -\lambda R + \sum_{x \in \mathcal{X}} \underbrace{\min_{m \in \mathcal{M}_x} [\boldsymbol{D}_{x,m} + \lambda \boldsymbol{R}_{x,m}]}_{\text{for fixed } (\lambda, x)} \right\}$$
- Fix $\lambda \geq 0$, $x \in \mathcal{X}$

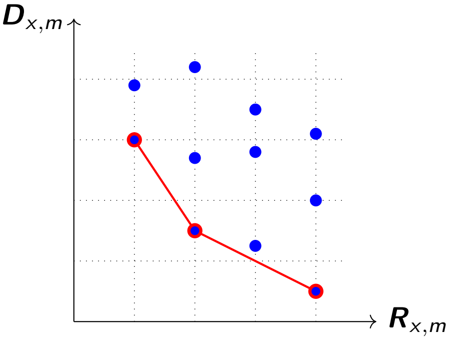

- Relevant points: $32{,}000^{10}$

Distortion-rate function: geometric solution via dual

- Dual: $$D^*(R) \;=\; \sup_{\lambda \geq 0} \Big\{ -\lambda R + \sum_{x \in \mathcal{X}} \underbrace{\min_{m \in \mathcal{M}_x} [\boldsymbol{D}_{x,m} + \lambda \boldsymbol{R}_{x,m}]}_{\text{for fixed } (\lambda, x)} \Big\}$$

- Fix $\lambda \geq 0$, $x \in \mathcal{X}$

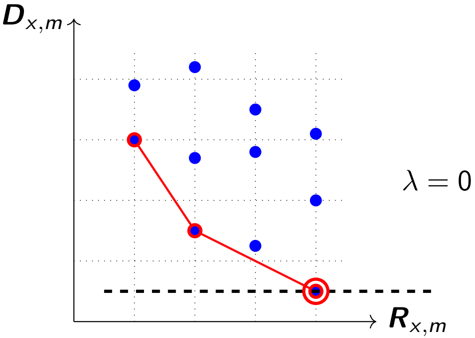

- Relevant points: $32{,}000^{10}$

Distortion-rate function: geometric solution via dual

- Dual: $$D^*(R) = \sup_{\lambda \geq 0} \Big\{ -\lambda R + \sum_{x \in \mathcal{X}} \underbrace{\min_{m \in \mathcal{M}_x} [\boldsymbol{D}_{x,m} + \lambda \boldsymbol{R}_{x,m}]}_{\text{for fixed } (\lambda, x)} \Big\}$$
- Fix $\lambda \geq 0$, $x \in \mathcal{X}$

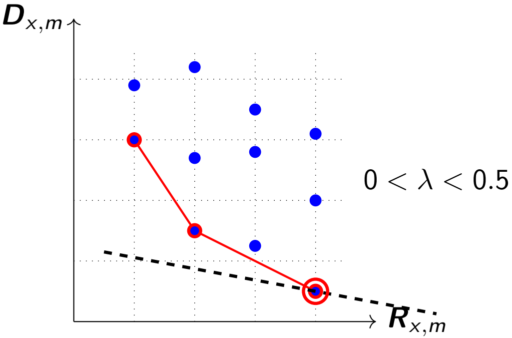

- Relevant points: $32{,}000^{10}$

Distortion-rate function: geometric solution via dual

- Dual: $$D^*(R) \;=\; \sup_{\lambda \geq 0} \Big\{ -\lambda R + \sum_{x \in \mathcal{X}} \underbrace{\min_{m \in \mathcal{M}_x} [\boldsymbol{D}_{x,m} + \lambda \boldsymbol{R}_{x,m}]}_{\text{for fixed } (\lambda, x)} \Big\}$$

- Fix $\lambda \geq 0$, $x \in \mathcal{X}$

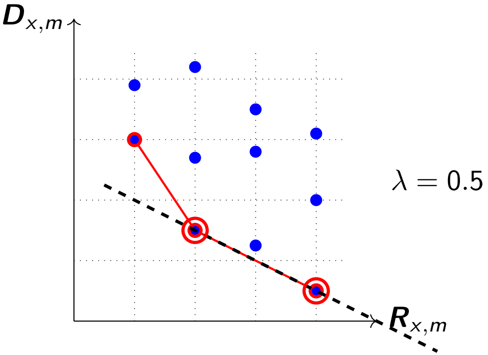

- Relevant points: $32{,}000^{10}$

Distortion-rate function: geometric solution via dual

- Dual: $$D^*(R) \; = \; \sup_{\lambda \geq 0} \Big\{ -\lambda R + \sum_{x \in \mathcal{X}} \underbrace{\min_{m \in \mathcal{M}_x} [\boldsymbol{D}_{x,m} + \lambda \boldsymbol{R}_{x,m}]}_{\text{for fixed } (\lambda, x)} \Big\}$$

- Fix $\lambda \geq 0$, $x \in \mathcal{X}$

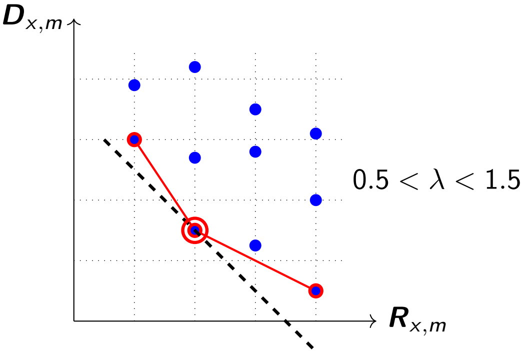

- Relevant points: $32{,}000^{10}$

Distortion-rate function: geometric solution via dual

- Dual: $$D^*(R) \;=\; \sup_{\lambda \geq 0} \left\{ -\lambda R + \sum_{x \in \mathcal{X}} \underbrace{\min_{m \in \mathcal{M}_x} \left[\boldsymbol{D}_{x,m} + \lambda \boldsymbol{R}_{x,m} \right]}_{\text{for fixed } (\lambda, x)} \right\}$$

- Fix $\lambda \geq 0$, $x \in \mathcal{X}$

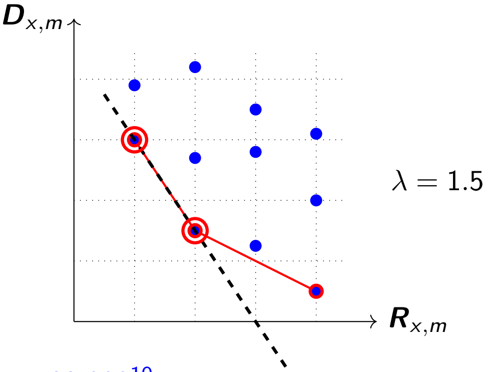

- Relevant points: $32{,}000^{10}$

Distortion-rate function: geometric solution via dual

- Dual: $$D^*(R) = \sup_{\lambda \geq 0} \left\{ -\lambda R + \sum_{x \in \mathcal{X}} \underbrace{\min_{m \in \mathcal{M}_x} [\boldsymbol{D}_{x,m} + \lambda \boldsymbol{R}_{x,m}]}_{\text{for fixed } (\lambda, x)} \right\}$$
- Fix $\lambda \geq 0$, $x \in \mathcal{X}$

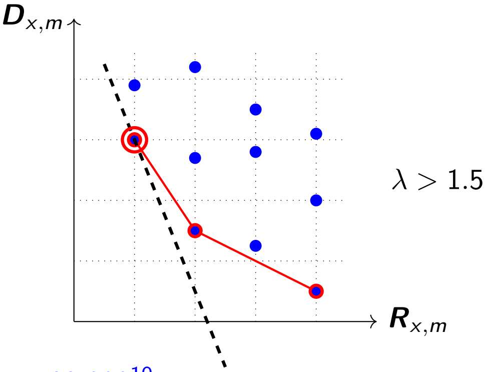

- Relevant points: $32{,}000^{10}$

Distortion-rate function: geometric solution via dual

- Dual: $$D^*(R) \;=\; \sup_{\lambda \geq 0} \left\{ -\lambda R + \sum_{x \in \mathcal{X}} \underbrace{\min_{m \in \mathcal{M}_x} \left[\boldsymbol{D}_{x,m} + \lambda \boldsymbol{R}_{x,m} \right]}_{\text{for fixed } (\lambda, x)} \right\}$$

- Fix $\lambda \geq 0$, $x \in \mathcal{X}$

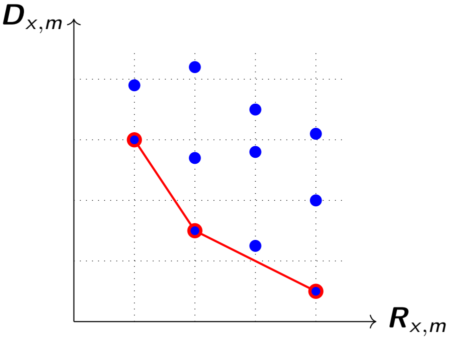

- Relevant points: $32{,}000^{10} \rightarrow 10$

Distortion-rate function: geometric solution via dual

- Dual: $$D^*(R) \;=\; \sup_{\lambda \geq 0} \Big\{ -\lambda R + \sum_{x \in \mathcal{X}} \underbrace{\min_{m \in \mathcal{M}_x} [\boldsymbol{D}_{x,m} + \lambda \boldsymbol{R}_{x,m}]}_{\text{for fixed } (\lambda, x)} \Big\}$$

- Fix $\lambda \geq 0$, $x \in \mathcal{X}$

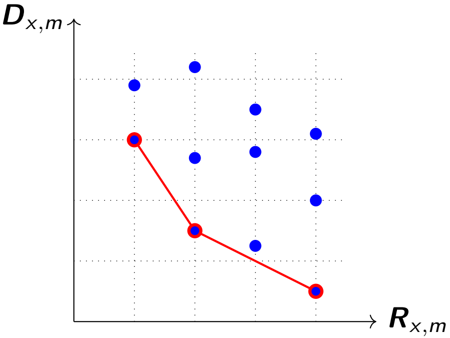

- Relevant points: $32{,}000^{10} \to 10$,
  only finitely many λ

Distortion-rate function: geometric solution via dual

- Dual: $$D^*(R) = \sup_{\lambda \geq 0} \Big\{ -\lambda R + \sum_{x \in \mathcal{X}} \underbrace{\min_{m \in \mathcal{M}_x} [\boldsymbol{D}_{x,m} + \lambda \boldsymbol{R}_{x,m}]}_{\text{for fixed } (\lambda, x)} \Big\}$$

- Fix $\lambda \geq 0$, $x \in \mathcal{X}$

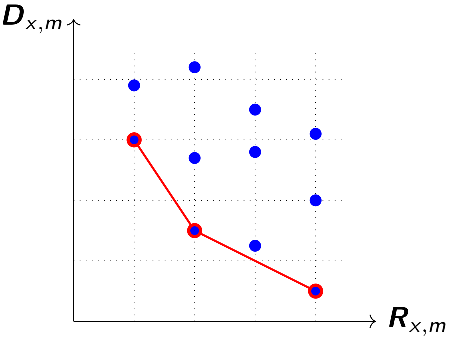

- Relevant points: $32{,}000^{10} \to 10$, $(2^{10} \to 10)$
  only finitely many λ

Distortion-rate function: geometric solution via dual

- Dual: $$D^*(R) \;=\; \sup_{\lambda \geq 0} \Big\{ -\lambda R + \sum_{x \in \mathcal{X}} \underbrace{\min_{m \in \mathcal{M}_x} [\boldsymbol{D}_{x,m} + \lambda \boldsymbol{R}_{x,m}]}_{\text{for fixed } (\lambda,x)} \Big\}$$
- Fix $\lambda \geq 0$, $x \in \mathcal{X}$

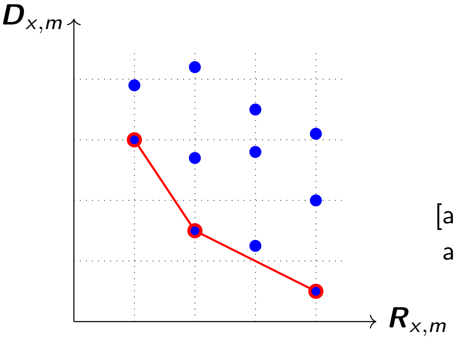

[apple $\mapsto$ app, ale, pe;
apple $\not\mapsto$ pale, red, lp]

- Relevant points: $32{,}000^{10} \to 10$, only finitely many λ

$(2^{10} \to 10)$

Experimental results

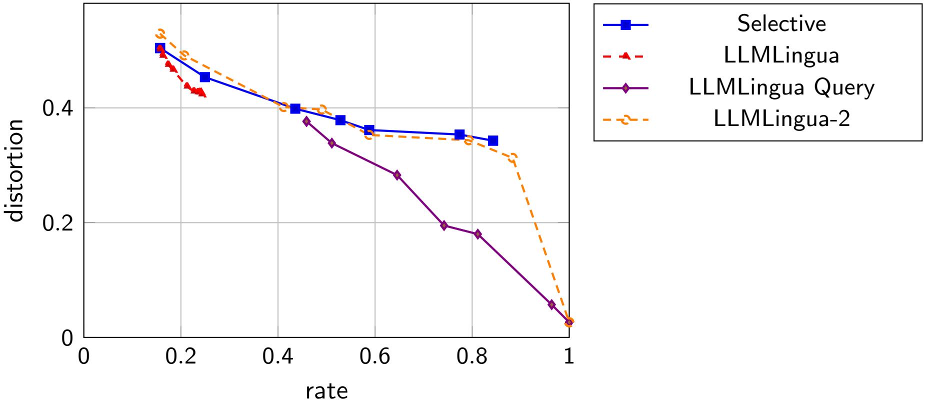

Experimental results

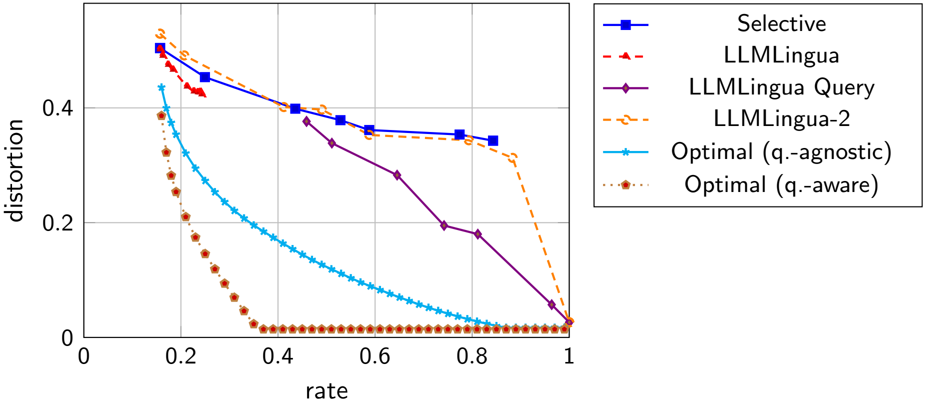

Experimental results

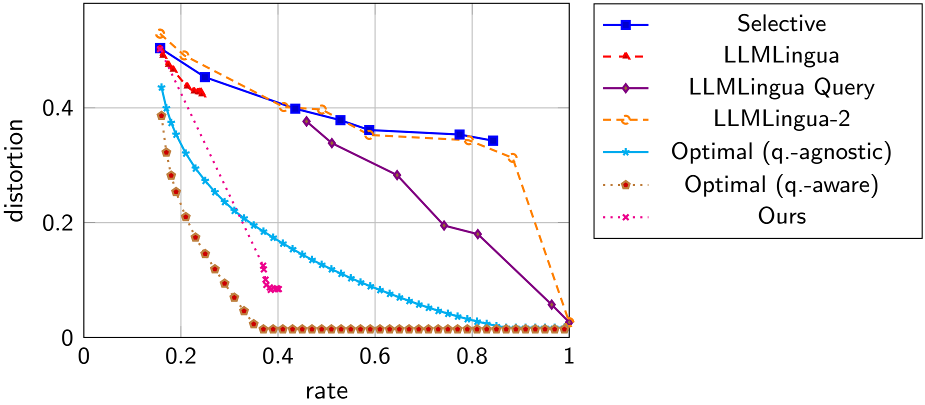

Experimental results

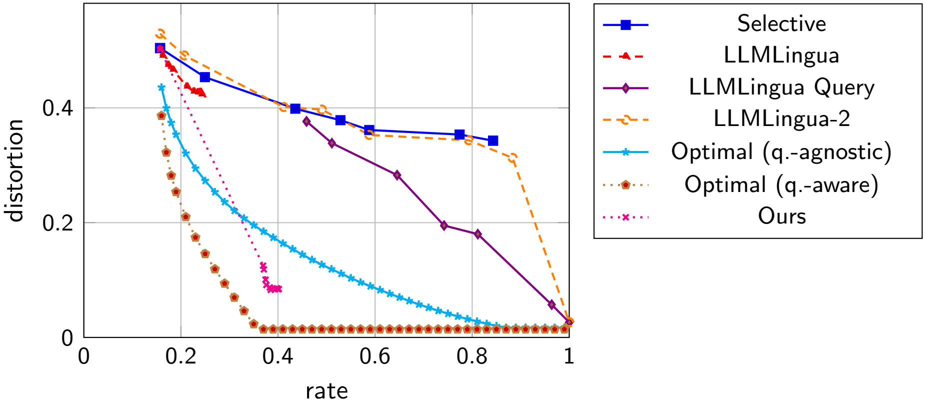

A.G.*, A.Nagle*, M.Bondaschi, M.Gastpar, A.V.Makkuva, H.Kim, "Fundamental Limits of Prompt Compression: A Rate-Distortion Framework for Black-Box Language Models."
— ICML 2024 Workshop on Theoretical Foundations of Foundation Models **[Oral]**
— under review at NeurIPS 2024

Segue to a contraction problem

- Optimization 101...

Segue to a contraction problem

- Optimization 101...thanks to a different problem:

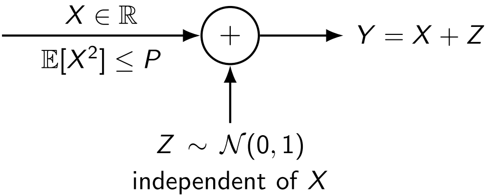

Segue to a contraction problem

- Optimization 101…thanks to a different problem:

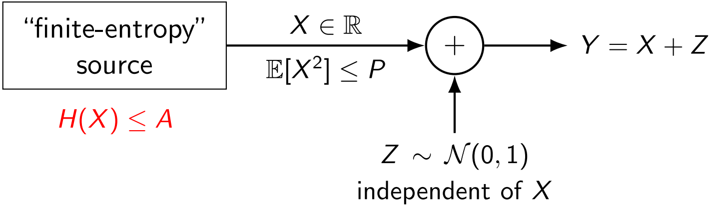

Outline

Input-entropy-constrained channel capacity

- A contraction problem in communication

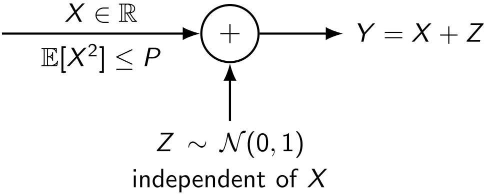

$$C_H(A, P) = \sup_{\substack{P_X:\ \mathbb{E}[X^2] \leq P \\ H(X) \leq A}} I(X; Y)$$

Input-entropy-constrained channel capacity

- A contraction problem in communication

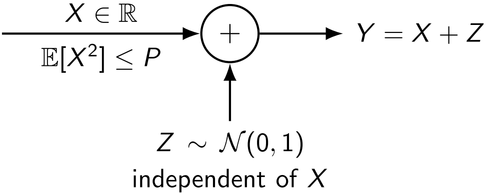

$$C_H(A, P) = \sup_{\substack{\mathsf{P}_X:\ \mathbb{E}[X^2] \leq P \\ H(X) \leq A}} I(X; Y)$$

- Cardinality bounds?

Input-entropy-constrained channel capacity

- A contraction problem in communication

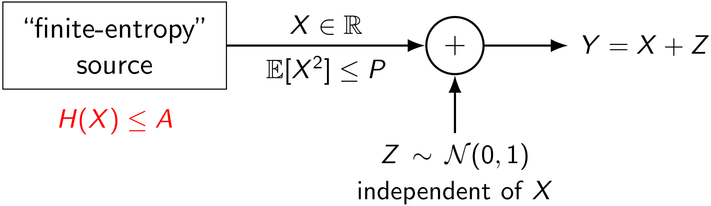

$$C_H(A, P) = \sup_{P_X:\ \substack{\mathbb{E}[X^2] \leq P \\ H(X) \leq A}} I(X; Y)$$

- Cardinality bounds? Finite support?

Input-entropy-constrained channel capacity

- A contraction problem in communication

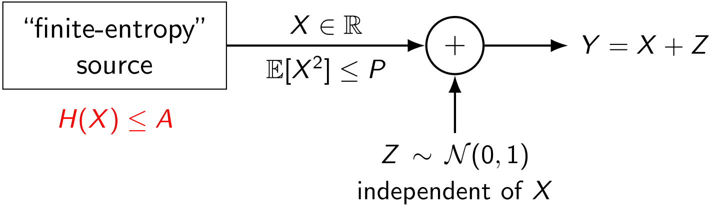

$$C_H(A,P) = \sup_{P_X:\ \substack{\mathbb{E}[X^2] \leq P \\ H(X) \leq A}} I(X;Y)$$

- Cardinality bounds? Finite support?
- A nontrivial upper bound better than

$$F_I(A,P) = \sup_{P_{WX}:\ \substack{\mathbb{E}[X^2] \leq P \\ I(W;X) \leq A}} I(W;Y) \qquad ?$$

Aside on data processing inequalities

Fix $\mathsf{P}_{Y|X}$

Aside on data processing inequalities

Fix $\mathsf{P}_{Y|X}$

- DPI: for any P_{WX}, $I(W;Y) \leq I(W;X)$

Aside on data processing inequalities

Fix $\mathsf{P}_{Y|X}$

- DPI: for any P_{WX}, $I(W;Y) \le I(W;X)$
- Data processing *function*: $F_I(t) = \sup_{\mathsf{P}_{WX}:\, I(W;X) \le t} I(W;Y)$

Aside on data processing inequalities

Fix $\mathsf{P}_{Y|X}$

- DPI: for any P_{WX}, $I(W;Y) \leq I(W;X)$
- Data processing *function*: $F_I(t) = \sup_{\mathsf{P}_{WX}:\, I(W;X) \leq t} I(W;Y)$

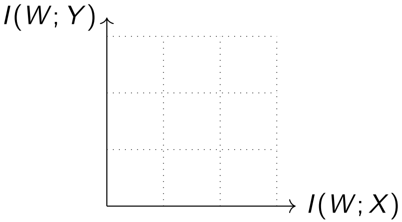

Aside on data processing inequalities

Fix $\mathsf{P}_{Y|X}$

- DPI: for any P_{WX}, $I(W;Y) \leq I(W;X)$
- Data processing *function*: $F_I(t) = \sup_{\mathsf{P}_{WX}:\, I(W;X) \leq t} I(W;Y)$

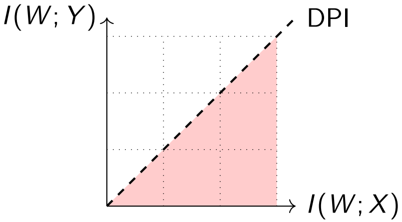

Aside on data processing inequalities

Fix $\mathrm{P}_{Y|X}$

- DPI: for any P_{WX}, $I(W;Y) \leq I(W;X)$
- Data processing *function*: $F_I(t) = \sup_{\mathrm{P}_{WX}:\, I(W;X) \leq t} I(W;Y)$

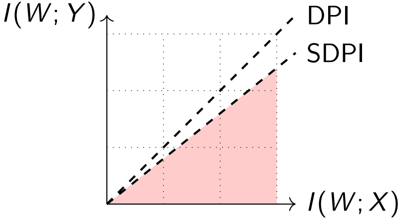

Aside on data processing inequalities

Fix $\mathsf{P}_{Y|X}$

- DPI: for any P_{WX}, $I(W;Y) \le I(W;X)$
- Data processing *function*: $F_I(t) = \sup_{\mathsf{P}_{WX}:\, I(W;X) \le t} I(W;Y)$

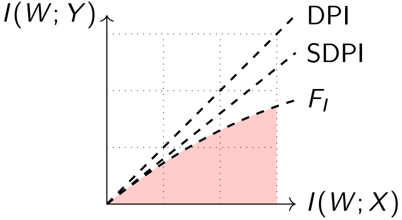

Aside on data processing inequalities

Fix $\mathrm{P}_{Y|X}$

- DPI: for any P_{WX}, $I(W;Y) \le I(W;X)$
- Data processing *function*: $F_I(t) = \sup_{\mathrm{P}_{WX}:\, I(W;X) \le t} I(W;Y)$

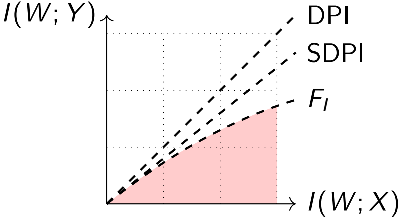

- Also DPI: for any $\mathrm{P}_X, \mathrm{Q}_X$, $D_f(\mathrm{Q}_Y \,\|\, \mathrm{P}_Y) \le D_f(\mathrm{Q}_X \,\|\, \mathrm{P}_X)$

Aside on data processing inequalities

Fix $\mathrm{P}_{Y|X}$

- DPI: for any P_{WX}, $I(W;Y) \le I(W;X)$

- Data processing *function*: $F_I(t) = \sup\limits_{\mathrm{P}_{WX}:\, I(W;X) \le t} I(W;Y)$

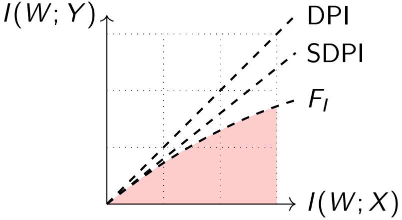

- Also DPI: for any $\mathrm{P}_X, \mathrm{Q}_X$, $D_f(\mathrm{Q}_Y \,\|\, \mathrm{P}_Y) \le D_f(\mathrm{Q}_X \,\|\, \mathrm{P}_X)$

$\mathrm{P}_Y = \mathrm{P}_X \circ \mathrm{P}_{Y|X}$

Aside on data processing inequalities

Fix $\mathrm{P}_{Y|X}$

- DPI: for any P_{WX}, $I(W;Y) \le I(W;X)$
- Data processing *function*: $F_I(t) = \sup_{\mathrm{P}_{WX}:\, I(W;X) \le t} I(W;Y)$

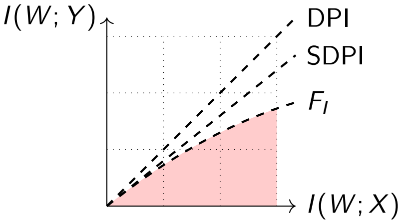

- Also DPI: for any $\mathrm{P}_X, \mathrm{Q}_X$, $D_f(\mathrm{Q}_Y \,\|\, \mathrm{P}_Y) \le D_f(\mathrm{Q}_X \,\|\, \mathrm{P}_X)$
- Natural analogue: $F_f(t) = \sup_{\mathrm{P}_X, \mathrm{Q}_X:\, D_f(\mathrm{Q}_X \,\|\, \mathrm{P}_X) \le t} D_f(\mathrm{Q}_Y \,\|\, \mathrm{P}_Y)$

Aside on data processing inequalities

Fix $\mathrm{P}_{Y|X}$

- DPI: for any P_{WX}, $I(W;Y) \leq I(W;X)$
- Data processing *function*: $F_I(t) = \sup_{\mathrm{P}_{WX}:\, I(W;X) \leq t} I(W;Y)$

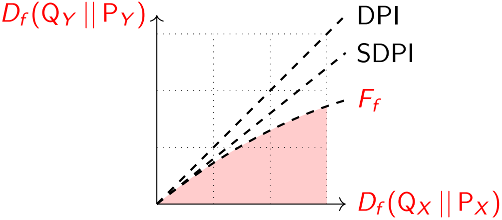

- Also DPI: for any $\mathrm{P}_X, \mathrm{Q}_X$, $D_f(\mathrm{Q}_Y \,||\, \mathrm{P}_Y) \leq D_f(\mathrm{Q}_X \,||\, \mathrm{P}_X)$
- Natural analogue: $F_f(t) = \sup_{\mathrm{P}_X, \mathrm{Q}_X:\, D_f(\mathrm{Q}_X \,||\, \mathrm{P}_X) \leq t} D_f(\mathrm{Q}_Y \,||\, \mathrm{P}_Y)$

Outline

Joint range of input and output divergences

Fix $\mathrm{P}_{Y|X}$ and f

Joint range of input and output divergences

Fix $\mathrm{P}_{Y|X}$ and f

- Define $F_f(t) = \sup_{\mathrm{P}_X, \mathrm{Q}_X \colon D_f(\mathrm{Q}_X \,\|\, \mathrm{P}_X) \leq t} D_f(\mathrm{Q}_Y \,\|\, \mathrm{P}_Y)$

Joint range of input and output divergences

Fix $\mathsf{P}_{Y|X}$ and f

- Define $F_f(t) = \sup_{\mathsf{P}_X, \mathsf{Q}_X :\, D_f(\mathsf{Q}_X \,||\, \mathsf{P}_X) \leq t} D_f(\mathsf{Q}_Y \,||\, \mathsf{P}_Y)$
- Upper boundary of $\mathcal{D}_f = \bigcup_{\mathsf{P}_X, \mathsf{Q}_X} \left\{ \big(D_f(\mathsf{Q}_X \,||\, \mathsf{P}_X), D_f(\mathsf{Q}_Y \,||\, \mathsf{P}_Y) \big) \right\}$

Joint range of input and output divergences

Fix $\mathsf{P}_{Y|X}$ and f

- Define $F_f(t) = \sup_{\mathsf{P}_X, \mathsf{Q}_X : D_f(\mathsf{Q}_X \,\|\, \mathsf{P}_X) \le t} D_f(\mathsf{Q}_Y \,\|\, \mathsf{P}_Y)$

- Upper boundary of $\mathcal{D}_f = \bigcup_{\mathsf{P}_X, \mathsf{Q}_X} \left\{ \left(D_f(\mathsf{Q}_X \,\|\, \mathsf{P}_X), D_f(\mathsf{Q}_Y \,\|\, \mathsf{P}_Y) \right) \right\}$

- Conjecture: $\mathcal{D}_f$ is convex ?

Joint range of input and output divergences

Fix $\mathrm{P}_{Y|X}$ and f

- Define $F_f(t) = \sup_{\mathrm{P}_X, \mathrm{Q}_X :\, D_f(\mathrm{Q}_X \,||\, \mathrm{P}_X) \le t} D_f(\mathrm{Q}_Y \,||\, \mathrm{P}_Y)$
- Upper boundary of $\mathcal{D}_f = \bigcup_{\mathrm{P}_X, \mathrm{Q}_X} \left\{ \left(D_f(\mathrm{Q}_X \,||\, \mathrm{P}_X), D_f(\mathrm{Q}_Y \,||\, \mathrm{P}_Y) \right) \right\}$
- Conjecture: $\mathcal{D}_f$ is convex ($\implies$ F_f is concave) ?

Joint range of input and output divergences

Fix $\mathsf{P}_{Y|X}$ and f

- Define $F_f(t) = \sup_{\mathsf{P}_X, \mathsf{Q}_X :\, D_f(\mathsf{Q}_X \,||\, \mathsf{P}_X) \le t} D_f(\mathsf{Q}_Y \,||\, \mathsf{P}_Y)$
- Upper boundary of $\mathcal{D}_f = \bigcup_{\mathsf{P}_X, \mathsf{Q}_X} \left\{ \left(D_f(\mathsf{Q}_X \,||\, \mathsf{P}_X), D_f(\mathsf{Q}_Y \,||\, \mathsf{P}_Y) \right) \right\}$
- Conjecture: $\mathcal{D}_f$ is convex ($\Longrightarrow$ F_f is concave) ?
- Facts:

Joint range of input and output divergences

Fix $\mathrm{P}_{Y|X}$ and f

- Define $F_f(t) = \sup_{\mathrm{P}_X, \mathrm{Q}_X : D_f(\mathrm{Q}_X \,||\, \mathrm{P}_X) \leq t} D_f(\mathrm{Q}_Y \,||\, \mathrm{P}_Y)$
- Upper boundary of $\mathcal{D}_f = \bigcup_{\mathrm{P}_X, \mathrm{Q}_X} \left\{ \left(D_f(\mathrm{Q}_X \,||\, \mathrm{P}_X), D_f(\mathrm{Q}_Y \,||\, \mathrm{P}_Y) \right) \right\}$
- Conjecture: $\mathcal{D}_f$ is convex ($\implies$ F_f is concave) ?
- Facts:
  - F_I is NOT necessarily concave (counter-example: $\mathrm{P}_{Y|X} = \mathrm{BEC}^3$)

Joint range of input and output divergences

Fix $\mathrm{P}_{Y|X}$ and f

- Define $F_f(t) = \sup_{\mathrm{P}_X, \mathrm{Q}_X :\, D_f(\mathrm{Q}_X \,\|\, \mathrm{P}_X) \le t} D_f(\mathrm{Q}_Y \,\|\, \mathrm{P}_Y)$
- Upper boundary of $\mathcal{D}_f = \bigcup_{\mathrm{P}_X, \mathrm{Q}_X} \left\{ \left(D_f(\mathrm{Q}_X \,\|\, \mathrm{P}_X), D_f(\mathrm{Q}_Y \,\|\, \mathrm{P}_Y) \right) \right\}$
- Conjecture: $\mathcal{D}_f$ is convex ($\Longrightarrow$ F_f is concave) ?
- Facts:
  - F_I is NOT necessarily concave (counter-example: $\mathrm{P}_{Y|X} = \mathrm{BEC}^3$)
  - Fix P_X, define $\tilde{F}_I(t, \mathrm{P}_X) = \sup_{\mathrm{P}_{W|X} :\, I(W;X) \le t} I(W;Y)$ and

Joint range of input and output divergences

Fix $\mathsf{P}_{Y|X}$ and f

- Define $F_f(t) = \sup_{\mathsf{P}_X, \mathsf{Q}_X \colon D_f(\mathsf{Q}_X \,||\, \mathsf{P}_X) \leq t} D_f(\mathsf{Q}_Y \,||\, \mathsf{P}_Y)$
- Upper boundary of $\mathcal{D}_f = \bigcup_{\mathsf{P}_X, \mathsf{Q}_X} \left\{ \left(D_f(\mathsf{Q}_X \,||\, \mathsf{P}_X), D_f(\mathsf{Q}_Y \,||\, \mathsf{P}_Y) \right) \right\}$
- Conjecture: $\mathcal{D}_f$ is convex ($\implies$ F_f is concave) ?
- Facts:
  - F_I is NOT necessarily concave (counter-example: $\mathsf{P}_{Y|X} = \mathrm{BEC}^3$)
  - Fix P_X, define $\tilde{F}_I(t, \mathsf{P}_X) = \sup_{\mathsf{P}_{W|X} \colon I(W;X) \leq t} I(W; Y)$ and

$$\tilde{F}_f(t, \mathsf{P}_X) = \sup_{\mathsf{Q}_X \colon D_f(\mathsf{Q}_X \,||\, \mathsf{P}_X) \leq t} D_f(\mathsf{Q}_Y \,||\, \mathsf{P}_Y),$$

Joint range of input and output divergences

Fix $\mathrm{P}_{Y|X}$ and f

- Define $F_f(t) = \sup_{\mathrm{P}_X, \mathrm{Q}_X :\, D_f(\mathrm{Q}_X \,\|\, \mathrm{P}_X) \le t} D_f(\mathrm{Q}_Y \,\|\, \mathrm{P}_Y)$
- Upper boundary of $\mathcal{D}_f = \bigcup_{\mathrm{P}_X, \mathrm{Q}_X} \left\{ \left(D_f(\mathrm{Q}_X \,\|\, \mathrm{P}_X), D_f(\mathrm{Q}_Y \,\|\, \mathrm{P}_Y) \right) \right\}$
- Conjecture: $\mathcal{D}_f$ is convex ($\Longrightarrow$ F_f is concave) ?
- Facts:
  - F_I is NOT necessarily concave (counter-example: $\mathrm{P}_{Y|X} = \mathrm{BEC}^3$)
  - Fix P_X, define $\tilde{F}_I(t, \mathrm{P}_X) = \sup_{\mathrm{P}_{W|X} :\, I(W;X) \le t} I(W;Y)$ and
  $$\tilde{F}_f(t, \mathrm{P}_X) = \sup_{\mathrm{Q}_X :\, D_f(\mathrm{Q}_X \,\|\, \mathrm{P}_X) \le t} D_f(\mathrm{Q}_Y \,\|\, \mathrm{P}_Y),$$
  then $\tilde{F}_I(\cdot, \mathrm{P}_X)$ is concave

Joint range of input and output divergences

Fix $\mathrm{P}_{Y|X}$ and f

- Define $F_f(t) = \sup_{\mathrm{P}_X, \mathrm{Q}_X :\, D_f(\mathrm{Q}_X \,||\, \mathrm{P}_X) \le t} D_f(\mathrm{Q}_Y \,||\, \mathrm{P}_Y)$
- Upper boundary of $\mathcal{D}_f = \bigcup_{\mathrm{P}_X, \mathrm{Q}_X} \left\{ \left(D_f(\mathrm{Q}_X \,||\, \mathrm{P}_X), D_f(\mathrm{Q}_Y \,||\, \mathrm{P}_Y) \right) \right\}$
- Conjecture: $\mathcal{D}_f$ is convex ($\implies$ F_f is concave) ?
- Facts:
  - F_I is NOT necessarily concave (counter-example: $\mathrm{P}_{Y|X} = \mathrm{BEC}^3$)
  - Fix P_X, define $\tilde{F}_I(t, \mathrm{P}_X) = \sup_{\mathrm{P}_{W|X} :\, I(W;X) \le t} I(W;Y)$ and

    $$\tilde{F}_f(t, \mathrm{P}_X) = \sup_{\mathrm{Q}_X :\, D_f(\mathrm{Q}_X \,||\, \mathrm{P}_X) \le t} D_f(\mathrm{Q}_Y \,||\, \mathrm{P}_Y),$$

    then $\tilde{F}_I(\cdot, \mathrm{P}_X)$ is concave; $\implies$ $\tilde{F}_f(\cdot, \mathrm{P}_X)$ is concave

Joint range of input and output divergences

Fix $\mathrm{P}_{Y|X}$ and f

- Define $F_f(t) = \sup_{\mathrm{P}_X, \mathrm{Q}_X \colon D_f(\mathrm{Q}_X \,||\, \mathrm{P}_X) \le t} D_f(\mathrm{Q}_Y \,||\, \mathrm{P}_Y)$
- Upper boundary of $\mathcal{D}_f = \bigcup_{\mathrm{P}_X, \mathrm{Q}_X} \left\{ \left(D_f(\mathrm{Q}_X \,||\, \mathrm{P}_X), D_f(\mathrm{Q}_Y \,||\, \mathrm{P}_Y) \right) \right\}$
- Conjecture: $\mathcal{D}_f$ is convex ($\implies$ F_f is concave) ?
- Facts:
  - F_I is NOT necessarily concave (counter-example: $\mathrm{P}_{Y|X} = \mathrm{BEC}^3$)
  - Fix P_X, define $\tilde{F}_I(t, \mathrm{P}_X) = \sup_{\mathrm{P}_{W|X} \colon I(W;X) \le t} I(W;Y)$ and

    $$\tilde{F}_f(t, \mathrm{P}_X) = \sup_{\mathrm{Q}_X \colon D_f(\mathrm{Q}_X \,||\, \mathrm{P}_X) \le t} D_f(\mathrm{Q}_Y \,||\, \mathrm{P}_Y),$$

    then $\tilde{F}_I(\cdot, \mathrm{P}_X)$ is concave; $\implies$ $\tilde{F}_f(\cdot, \mathrm{P}_X)$ is concave
  - For any f, g, $\bigcup_{\mathrm{P}_X, \mathrm{Q}_X} \left\{ \left(D_f(\mathrm{Q}_X \,||\, \mathrm{P}_X), D_g(\mathrm{Q}_X \,||\, \mathrm{P}_X) \right) \right\}$ is convex

Outline

In closing. . .

In closing. . .

- Three problems ($1\times$ compression, $2\times$ contraction):

In closing...

- Three problems ($1\times$ compression, $2\times$ contraction):
  - Prompt compression for LLMs

In closing...

- Three problems ($1\times$ compression, $2\times$ contraction):
  - Prompt compression for LLMs
  - Entropy-constrained capacity

In closing...

- Three problems ($1\times$ compression, $2\times$ contraction):
  - Prompt compression for LLMs
  - Entropy-constrained capacity
  - Joint range of divergences

In closing...

- Three problems ($1\times$ compression, $2\times$ contraction):
  - Prompt compression for LLMs
  - Entropy-constrained capacity
  - Joint range of divergences
- Two more:

In closing...

- Three problems (1$\times$ compression, 2$\times$ contraction):
  - Prompt compression for LLMs
  - Entropy-constrained capacity
  - Joint range of divergences
- Two more:
  - Guesswork

In closing...

- Three problems ($1\times$ compression, $2\times$ contraction):
  - Prompt compression for LLMs
  - Entropy-constrained capacity
  - Joint range of divergences
- Two more:
  - Guesswork
  - Distributed hypothesis testing

In closing. . .

- Three problems ($1\times$ compression, $2\times$ contraction):
  - Prompt compression for LLMs
  - Entropy-constrained capacity
  - Joint range of divergences
- Two more (method of types + optimization):
  - Guesswork
  - Distributed hypothesis testing

In closing...

- Three problems ($1\times$ compression, $2\times$ contraction):
  - Prompt compression for LLMs
  - Entropy-constrained capacity
  - Joint range of divergences
- Two more (method of types + optimization):
  - Guesswork
  - Distributed hypothesis testing $\longrightarrow$ compression + contraction

In closing...

- Three problems ($1\times$ compression, $2\times$ contraction):
  - Prompt compression for LLMs
  - Entropy-constrained capacity
  - Joint range of divergences
- Two more (method of types + optimization):
  - Guesswork
  - Distributed hypothesis testing $\longrightarrow$ compression + contraction
- All thoughts welcome

In closing...

- Three problems ($1\times$ compression, $2\times$ contraction):
  - Prompt compression for LLMs
  - Entropy-constrained capacity
  - Joint range of divergences
- Two more (method of types + optimization):
  - Guesswork
  - Distributed hypothesis testing $\longrightarrow$ compression + contraction
- All thoughts welcome

Thank you!